

# Fast, Cheap, DIY Observability with Open Source Analytics and Visualization

Josh Lee, Altinity



**devm.io NEW YORK WEEK**



Josh Lee

Open Source Advocate  
*Altinity*

*ClickHouse® is a registered trademark of ClickHouse, Inc.  
Altinity is not affiliated with or associated with ClickHouse, Inc.  
We are but humble open source contributors*

Observability is our  
ability to understand  
a system from its  
outputs alone









# A Typical Request Log




```
2024-07-01 09:35:34 GET /home 200 ...
```

# Adding Duration



```
2024-07-01 09:35:34 231ms GET /home 200
```



## Overview



## Latency



# Back to our log...

```
2024-07-01 09:35:34 231ms GET /home 200
```

# Back to our log...

```
Request:123 2024-07-01 09:35:34 231ms GET /home  
200
```

# Connecting the trace:

```
Trace:4ea3 Span:123 2024-07-01 09:35:34 231ms  
GET /home 200
```

```
Trace:4ea3 Span:456 ParentSpan:123 2024-07-01  
09:35:34 201ms GET /api/users 201
```

# ▼ frontend: HTTP GET ca28836


Find...

⊙ ^ v ×

⌘

Trace Timeline

Start May 31 2023, 11:36:59.537 Duration 24.95ms Services 4 Depth 7 Total Spans 17



Service & Operation	0μs	6.24ms	12.48ms	18.71ms
frontend HTTP GET				
frontend grpc.oteldemo.RecommendationService/List...				17.69ms
▼ recommendationservice /oteldemo.Recomm...				14.12ms
▼ recommendationservice get_product_list				13.61ms
▼ recommendationservice /oteldemo...				11.54ms   recommendationservice::oteldemo.FeatureFlagService/GetF...
▼ featureflagservice /oteldemo.Fe...			2.23ms	
featureflagservice featurerefl...			1.93ms	
▼ recommendationservice /oteldemo...				1.62ms
productcatalogservice otelde...				29μs
frontend grpc.oteldemo.ProductCatalogService/GetPr...				5.48ms
productcatalogservice oteldemo.ProductCat...				27μs
frontend grpc.oteldemo.ProductCatalogService/GetPr...				4.42ms
productcatalogservice oteldemo.ProductCat...				9μs
frontend grpc.oteldemo.ProductCatalogService/GetPr...				5.44ms
productcatalogservice oteldemo.ProductCat...				21μs
frontend grpc.oteldemo.ProductCatalogService/GetPr...				3.16ms
productcatalogservice oteldemo.ProductCat...				20μs



# Observability is not any one signal...

## Metrics

Aggregable

*Is there a problem?*

## Traces

Request-Scoped

*Where is the problem?*

## Logs


Verbose, time-stamped records

*What is the problem?*

```
$ vmstat -n 2 10
```

procs	-----memory-----	---swap--	-----io----	-system--	-----cpu-----												
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st	
0	0	343296	21690808	2290104	6897160	0	0	0	3	187	0	2	4	2	94	0	0
0	0	343296	21690800	2290104	6897160	0	0	0	0	60	2989	7688	2	1	97	0	0
0	0	343296	21690140	2290104	6897164	0	0	0	0	72	4704	13677	3	2	95	0	0
0	0	343296	21689888	2290104	6897164	0	0	0	0	14	3132	9364	2	1	97	0	0
0	0	343296	21690220	2290104	6897168	0	0	0	0	86	3014	7995	1	1	97	0	0
0	0	343296	21690448	2290104	6897176	0	0	0	0	20	2660	7297	1	1	98	0	0
0	0	343296	21690268	2290104	6897176	0	0	0	0	12	2695	7222	1	1	98	0	0
1	0	343296	21690196	2290104	6897180	0	0	0	0	80	3641	10419	2	1	97	0	0
0	0	343296	21689696	2290104	6897180	0	0	0	0	14	4108	12605	3	2	95	0	0
0	0	343296	21689900	2290104	6897184	0	0	0	0	60	2688	7270	2	1	97	0	0

# A complete observability solution



# Introducing ClickHouse

- SQL-compatible
- Massively scalable
- Really, really fast

# Telemetry is WORM

Write-Once, Read-Many




# Telemetry is WORM

Write-Once, Read-Maybe


# B-Trees: Optimized for Reads

# Log-Structured Merge Trees: Optimized for ingestion




and periodically flushed to disk...forming a set of small, sorted files.

# Log-Structured Merge Trees: Background compaction





Compaction continues creating fewer, larger and larger files

Part



Rewritten, Bigger Part


Part



Update and delete also rewrite parts



Unmerged,  
freshly  
inserted  
part



Fully  
merged  
part

Query efficiency

ClickHouse for Observability

# How does this help?

- Fast writes
- Time-friendly
- Easy cleanup
- Cost-effective


ClickHouse for Observability

# Integrations

- Grafana Datasource Plugin
- Jaeger w/ ClickHouse backend
- Kafka table engine

ClickHouse for Observability

# Integrations via OpenTelemetry



```
CREATE TABLE default.ontime_ref(  
  `Year` UInt16,  
  `Quarter` UInt8,  
  `Month` UInt8,  
  `FlightDate` Date,  
  `Carrier` LowCardinality(String) ,  
  . . .  
)
```

```
ENGINE = MergeTree
```

```
PARTITION BY Year
```

```
ORDER BY (Carrier, FlightDate)
```

LZ4 compression

Dictionary encoding

Engine for fast  
analytics


How to partition data

How to sort rows



## PostgreSQL, MySQL


Read all columns in row




Rows compressed minimally or not at all

## ClickHouse

Read only selected columns




Columns highly compressed




Ok, back to our  
simple o11y demo...

```
$ vmstat -n 2 10
```

procs		-----memory-----				---swap--		-----io----		-system--		-----cpu-----					
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st	
0	0	343296	21690808	2290104	6897160	0	0	0	3	187	0	2	4	2	94	0	0
0	0	343296	21690800	2290104	6897160	0	0	0	0	60	2989	7688	2	1	97	0	0
0	0	343296	21690140	2290104	6897164	0	0	0	0	72	4704	13677	3	2	95	0	0
0	0	343296	21689888	2290104	6897164	0	0	0	0	14	3132	9364	2	1	97	0	0
0	0	343296	21690220	2290104	6897168	0	0	0	0	86	3014	7995	1	1	97	0	0
0	0	343296	21690448	2290104	6897176	0	0	0	0	20	2660	7297	1	1	98	0	0
0	0	343296	21690268	2290104	6897176	0	0	0	0	12	2695	7222	1	1	98	0	0
1	0	343296	21690196	2290104	6897180	0	0	0	0	80	3641	10419	2	1	97	0	0
0	0	343296	21689696	2290104	6897180	0	0	0	0	14	4108	12605	3	2	95	0	0
0	0	343296	21689900	2290104	6897184	0	0	0	0	60	2688	7270	2	1	97	0	0




# A complete observability solution



# Sooo...How do we ingest vmstat data and display it?

```
$ vmstat 1 -n
```

procs		-----memory-----				---swap--		-----io----		-system--		-----cpu-----				
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st
0	0	166912	2645740	36792	3360652	0	0	3	101	1	1	2	1	98	0	0
1	0	166912	2645360	36792	3360652	0	0	0	0	1182	3986	7	1	93	0	0



# Step 1: Generate vmstat data

```
#!/usr/bin/env python3
import datetime, json, socket, subprocess
host = socket.gethostname()
with subprocess.Popen(['vmstat', '-n', '1'], stdout=subprocess.PIPE) as proc:
    proc.stdout.readline() # discard first line
    header_names = proc.stdout.readline().decode().split()
    values = proc.stdout.readline().decode()
    while values != '' and proc.poll() is None:
        dict = {}
        dict['timestamp'] = datetime.datetime.now().strftime("%Y-%m-%d %H:%M:%S")
        dict['host'] = host
        for (header, value) in zip(header_names, values.split()):
            dict[header] = int(value)
        print(json.dumps(dict), flush=True)
        values = proc.stdout.readline().decode()
```



## Here's the output

```
{"timestamp": "2024-01-22 18:13:16", "host": "logos3", "r": 0, "b": 0, "swpd": 166912, "free": 2523688, "buff": 41412, "cache": 3408292, "si": 0, "so": 0, "bi": 3, "bo": 101, "in": 1, "cs": 0, "us": 2, "sy": 1, "id": 98, "wa": 0, "st": 0}
```

```
{"timestamp": "2024-01-22 18:13:17", "host": "logos3", "r": 0, "b": 0, "swpd": 166912, "free": 2523696, "buff": 41412, "cache": 3408316, "si": 0, "so": 0, "bi": 0, "bo": 216, "in": 1214, "cs": 4320, "us": 1, "sy": 1, "id": 98, "wa": 0, "st": 0}
```

```
{"timestamp": "2024-01-22 18:13:18", "host": "logos3", "r": 0, "b": 0, "swpd": 166912, "free": 2527120, "buff": 41412, "cache": 3408572, "si": 0, "so": 0, "bi": 0, "bo": 0, "in": 1172, "cs": 4162, "us": 2, "sy": 1, "id": 98, "wa": 0, "st": 0}
```

## Step 2: Design a ClickHouse table to hold data

```
CREATE TABLE monitoring.vmstat (  
  timestamp DateTime,  
  day UInt32 default toYYYYMMDD(timestamp),  
  host String,  
  r UInt64, b UInt64, -- procs  
  swpd UInt64, free UInt64, buff UInt64, cache UInt64, -- memory  
  si UInt64, so UInt64, -- swap  
  bi UInt64, bo UInt64, -- io  
  in UInt64, cs UInt64, -- system  
  us UInt64, sy UInt64, id UInt64, wa UInt64, st UInt64 -- cpu  
) ENGINE=MergeTree  
PARTITION BY day  
ORDER BY (host, timestamp)
```

The diagram illustrates the mapping of fields in the ClickHouse table definition to Dimensions and Measurements. A yellow box labeled "Dimensions" has arrows pointing to the fields `timestamp`, `day`, and `host`. Another yellow box labeled "Measurements" has arrows pointing to the fields `sy`, `id`, and `wa`.

## Step 3: Load data into ClickHouse


```
INSERT INTO vmstat Format JSONEachRow
```

```
INSERT='INSERT%20INTO%20vmstat%20Format%20JSONEachRow'  
cat vmstat.dat | curl -X POST --data-binary @- \  
  "http://logos3:8123/?database=monitoring&query=${INSERT}"
```

(Or a Python script)

Host Name logos2


Average CPU usage by host



CPU Usage: logos2



Memory Usage: logos2



## Step 5: Go crazy!


```
SELECT host, count() AS loaded_minutes
FROM (
    SELECT
        toStartOfMinute(timestamp) AS minute, host, avg(100 - id) AS load
    FROM monitoring.vmstat
    WHERE timestamp > (now() - toIntervalDay(1))
    GROUP BY minute, host HAVING load > 25
)
GROUP BY host ORDER BY loaded_minutes DESC
```

host	loaded_minutes
logos3	6
logos2	5

2 hosts had > 25% load for at least a minute in the last 24 hours

Host Name logos2 ▾

Average CPU usage by host



CPU Usage: logos2




stress -m 8 --vm-bytes 4G

Memory Usage: logos2



# Example #2: Monitoring Proxmox Hosts

# Monitoring Proxmox Hosts





source	key
Attributes	host
Attributes	instance
Attributes	nodename
Attributes	object
Attributes	type
Attributes	vmid
ResourceAttributes	host name

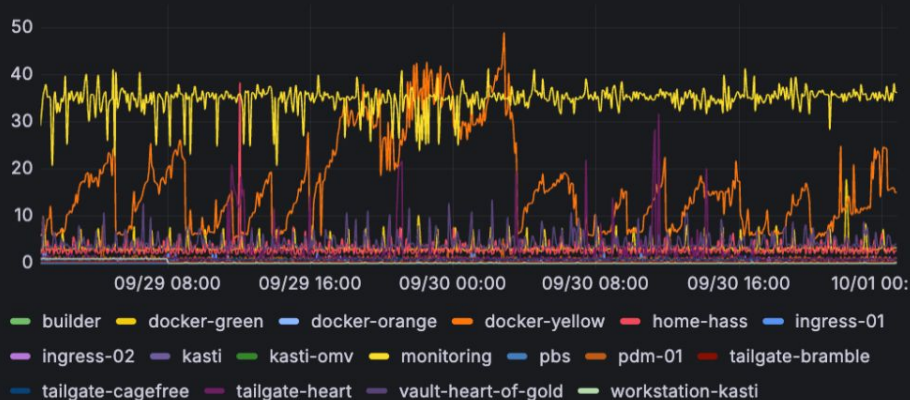
host	node
monitoring	ca
tailgate-cagefree	ca
home-hass	he
ingress-01	he
kasti	he
tailgate-heart	he
vault-heart-of-gold	he

node ↑
cagefree
heart-of-gold
pve-01
pve-02


source	key
otel_metrics_gauge	ballooninfo_actual
otel_metrics_gauge	ballooninfo_free_mem
otel_metrics_gauge	ballooninfo_last_update
otel_metrics_gauge	ballooninfo_major_page_fa
otel_metrics_gauge	ballooninfo_max_mem
otel_metrics_gauge	ballooninfo_mem_swapper
otel_metrics_gauge	ballooninfo_mem_swapper

## Overview

### CPU by VM



### Memory by VM



Data source



\${Datasource}



Query options

MD = auto = 495

Interval = 5m

Query inspector



A

(\${Datasource})



Query Settings

SQL Editor

Run Query

```
SELECT
  t,
  groupArray((host, cpu_percent)) as groupArr
FROM (
  SELECT
    $timeSeries as t,
    Attributes['host'] as host,
    max(Value * 100) as cpu_percent
  FROM metrics.otel_metrics_gauge
  WHERE MetricName = 'system_cpu'
    AND $timeFilter
    AND nodename IN($Node)
  GROUP BY t, host
  ORDER BY t, host
)
GROUP BY t
ORDER BY t
```

Extrapolation



Step



Resolution

1/1




Round

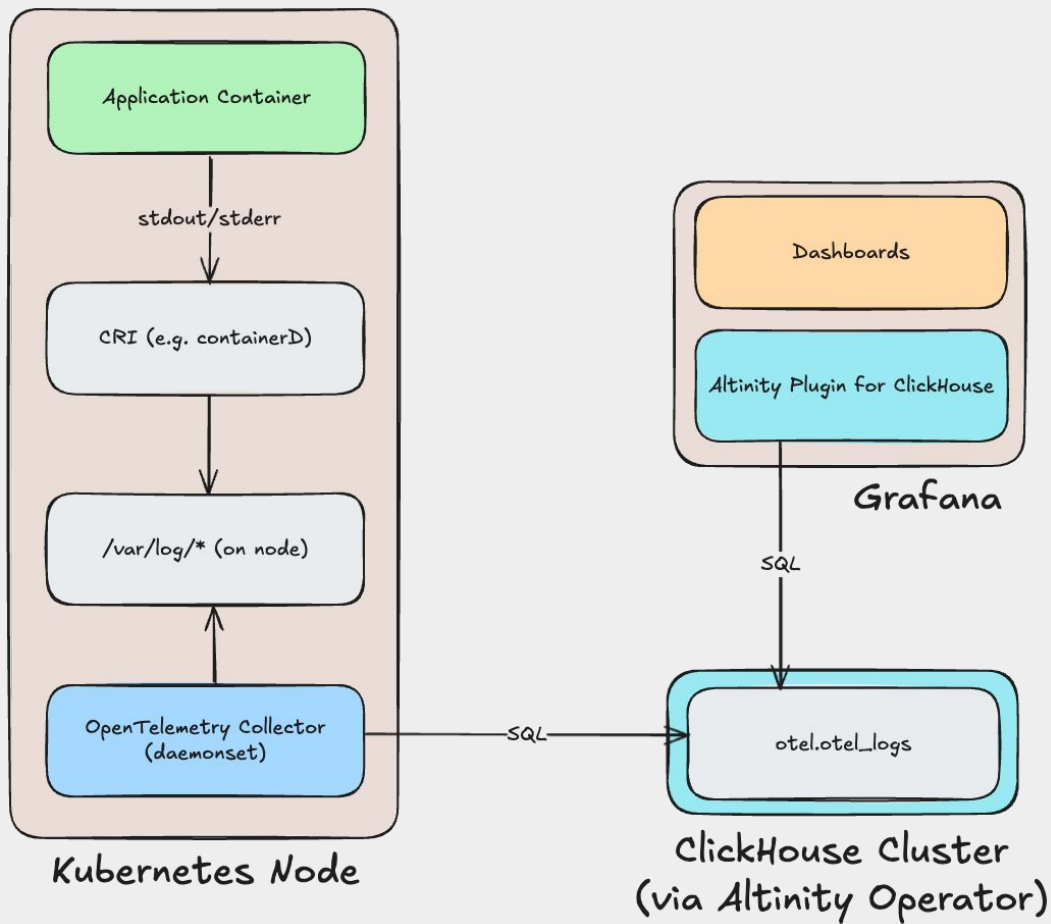


0s


# Example #3: Ingesting K8s Logs w/ OpenTelemetry + ClickHouse



OpenTelemetry Collector Pipeline



Finally... scaling for production



Cheap, fast, AND good?!



# What have we learned?

- Keep it simple
- Start with what you've got
- OpenTelemetry is a swiss-army knife
- ClickHouse is pretty cool

# Thank you and happy querying!

Josh Lee - Altinity



Connect with me



Resources & slides