

(Open Source) Observability from Scratch

Josh Lee • Altinity • May 1 2025 • OpenSearchCon Europe



Josh Lee
Open Source Advocate
Altinity

*Altinity® is a Registered Trademark of Altinity, Inc. ClickHouse® is a registered trademark of ClickHouse, Inc.;
Altinity is not affiliated with or associated with ClickHouse, Inc.
We are but humble open source contributors*

Observability is our
ability to understand a
system from its
outputs alone





“There are only two
signals: metrics and
(structured) logs”

— paraphrased from Charity Majors, Honeycomb




A Typical Request Log

```
2024-07-01 09:35:34 GET /home 200 ...
```



Adding Duration

```
2024-07-01 09:35:34 231ms GET /home 200
```


Overview



Latency



Back to our log...

```
2024-07-01 09:35:34 231ms GET /home 200
```

Back to our log...

```
Request:123 2024-07-01 09:35:34 231ms GET /home  
200
```


Connecting the trace:

```
Trace:4ea3 Span:123 2024-07-01 09:35:34 231ms GET  
/home 200
```

```
Trace:4ea3 Span:456 ParentSpan:123 2024-07-01  
09:35:34 201ms GET /api/users 201
```

▼ frontend: HTTP GET ca28836

Find...

◉ ^ v ×

⌘

Trace Timeline


Start May 31 2023, 11:36:59.537

Duration 24.95ms

Services 4

Depth 7

Total Spans 17



Name & Operation	0µs	6.24ms	12.48ms	18.71ms	
frontend HTTP GET	<div></div>				
frontend grpc.oteldemo.RecommendationService/List...	<div></div>				17.69ms
▼ recommendationservice /oteldemo.Recomm...	<div></div>				14.12ms
▼ recommendationservice get_product_list	<div></div>				13.61ms
▼ recommendationservice /oteldemo...	<div></div>				11.54ms recommendationservice::/oteldemo.FeatureFlagService/GetF...
▼ featureflagservice /oteldemo.Fe...	<div></div>				2.23ms
featureflagservice featurefl...	<div></div>				1.93ms
▼ recommendationservice /oteldemo...	<div></div>				1.62ms
productcatalogservice otelde...	<div></div>				29µs
frontend grpc.oteldemo.ProductCatalogService/GetPr...	<div></div>				5.48ms
productcatalogservice oteldemo.ProductCat...	<div></div>				27µs
frontend grpc.oteldemo.ProductCatalogService/GetPr...	<div></div>				4.42ms
productcatalogservice oteldemo.ProductCat...	<div></div>				9µs
frontend grpc.oteldemo.ProductCatalogService/GetPr...	<div></div>				5.44ms
productcatalogservice oteldemo.ProductCat...	<div></div>				21µs
frontend grpc.oteldemo.ProductCatalogService/GetPr...	<div></div>				3.16ms
productcatalogservice oteldemo.ProductCat...	<div></div>				20µs

Observability is not any one signal...

Metrics

Aggregable

Is there a problem?

Traces

Request-Scoped

Where is the problem?

Logs

Verbose, time-stamped records

What is the problem?

Distributed Tracing is the *"Killer App"*

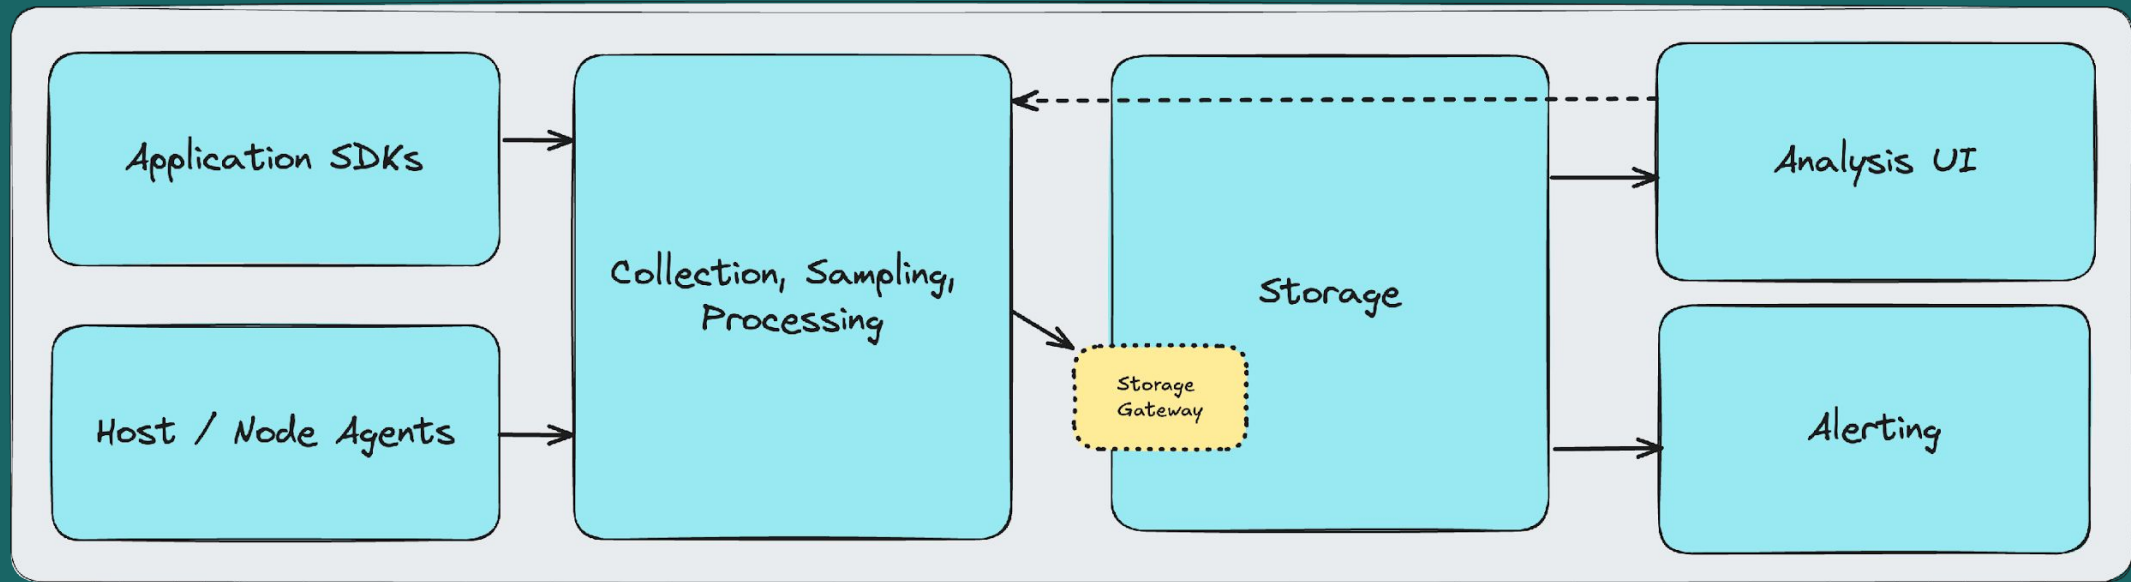
Understand
complete request
flows

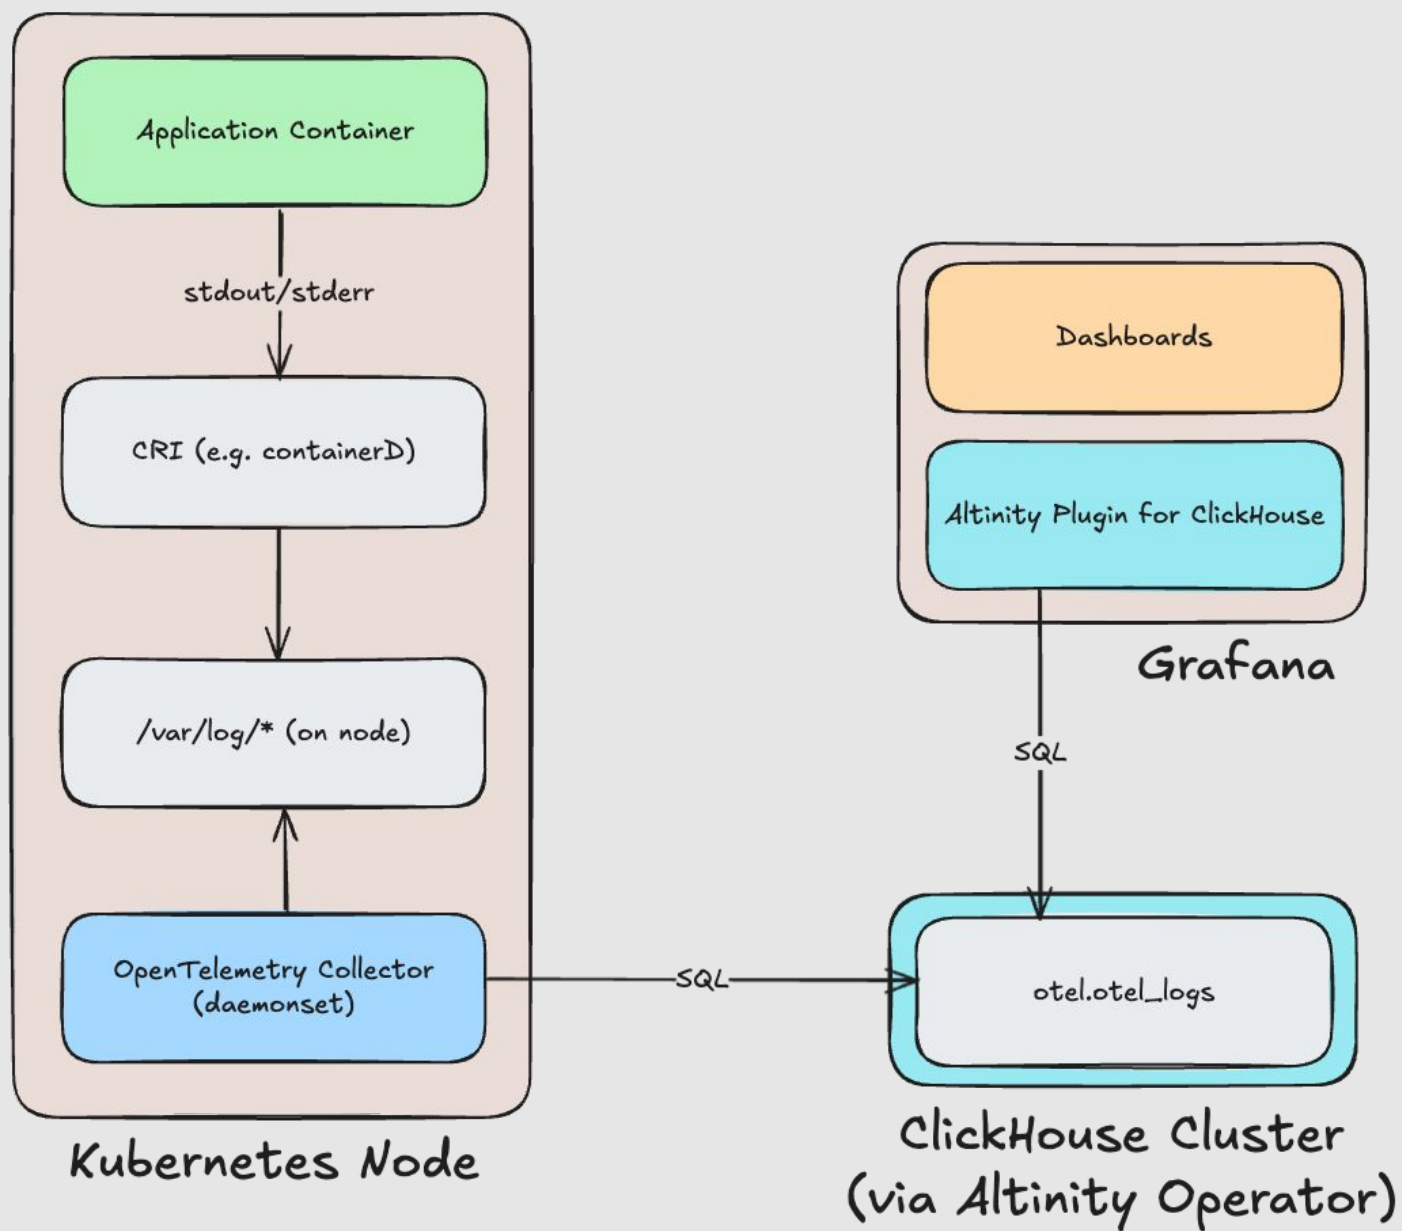
Create a
real-time map of
system topology
and
dependencies

Derive metrics
from the richness
of trace
metadata

Enrich logs and
metrics with
context

A complete observability solution






Introducing ClickHouse

SQL-compatible
Massively scalable
Really, really fast
Apache 2.0


Telemetry is WORM

Write-Once, Read-Many


Part




Part



Rewritten, Bigger Part





ClickHouse for Observability

ClickHouse for Observability

How does this help?

- Fast writes
- Time-friendly
- Easy cleanup
- Cost-effective

ClickHouse for Observability

Integrations

- Grafana Datasource Plugin
- Jaeger w/ ClickHouse backend
- qryn (Prometheus, Loki, Tempo, +more)
- Kafka table engine
- OpenTelemetry Exporter

Observability for ClickHouse


```
SELECT
  query_id,
  query_duration_ms,
  user,
  left(query, 20),
  written_bytes
FROM system.query_log
ORDER BY event_date DESC
LIMIT 1
```

Query id: e0214d77-e6b8-4e94-b34b-96dd13432364

	query_id	query_duration_ms	user	left(query, 20)	written_bytes
1.	605f93cf-5586-435b-8ffe-988b11cbc539	58	default	INSERT INTO otel_tra	726906

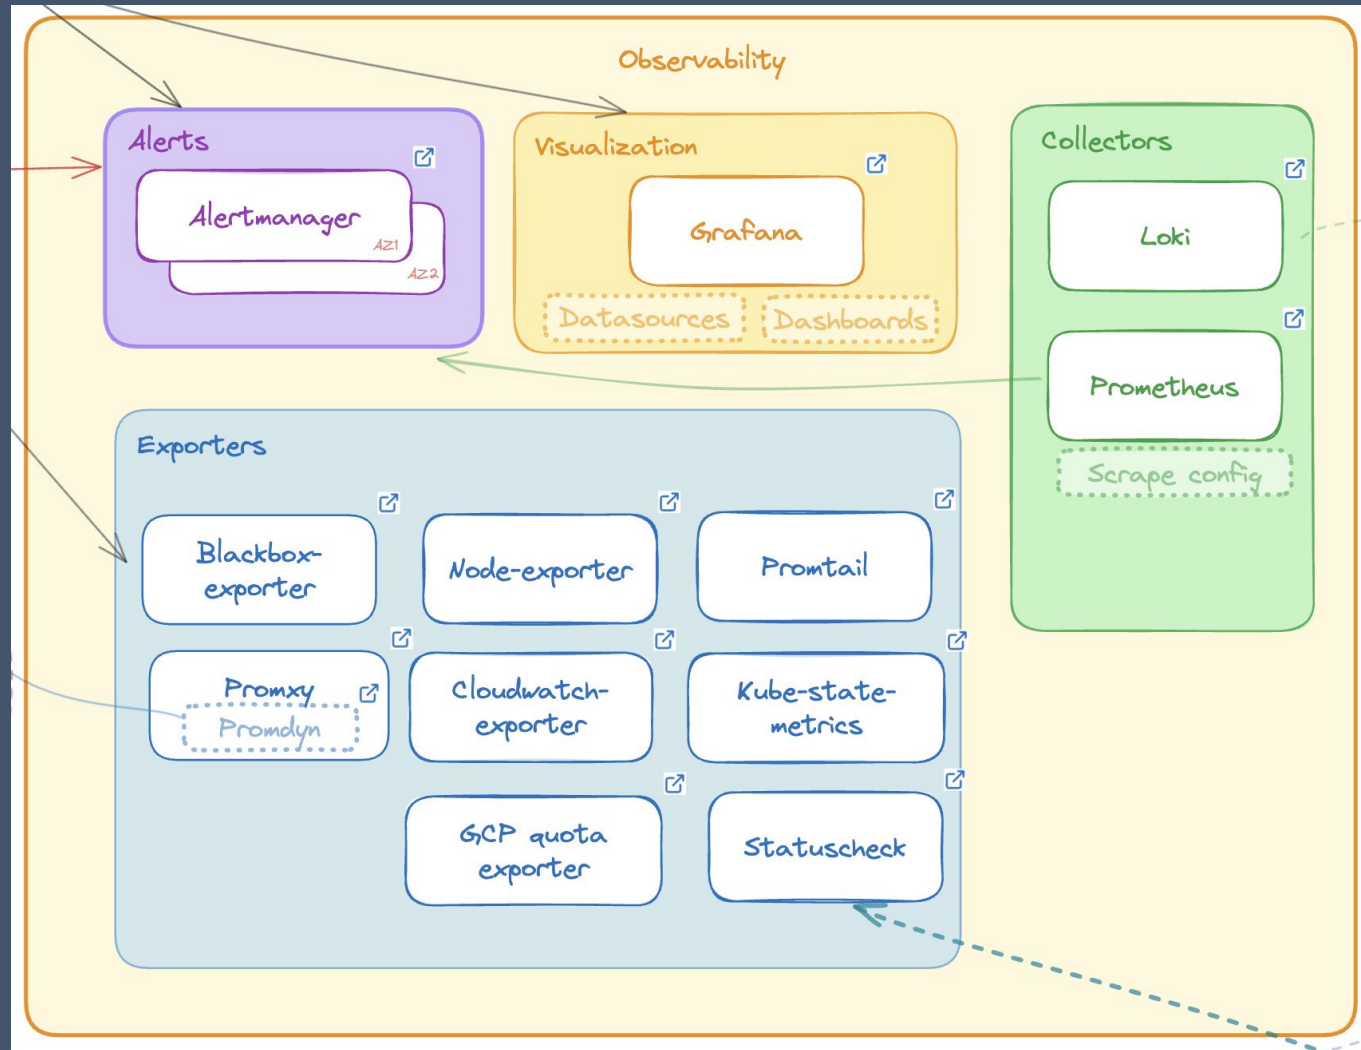
1 row in set. Elapsed: 0.017 sec. Processed 12.15 thousand rows, 7.99 MB (725.40 thousand rows/s., 477.05 MB/s.)
Peak memory usage: 16.69 MiB.

Observability for ClickHouse

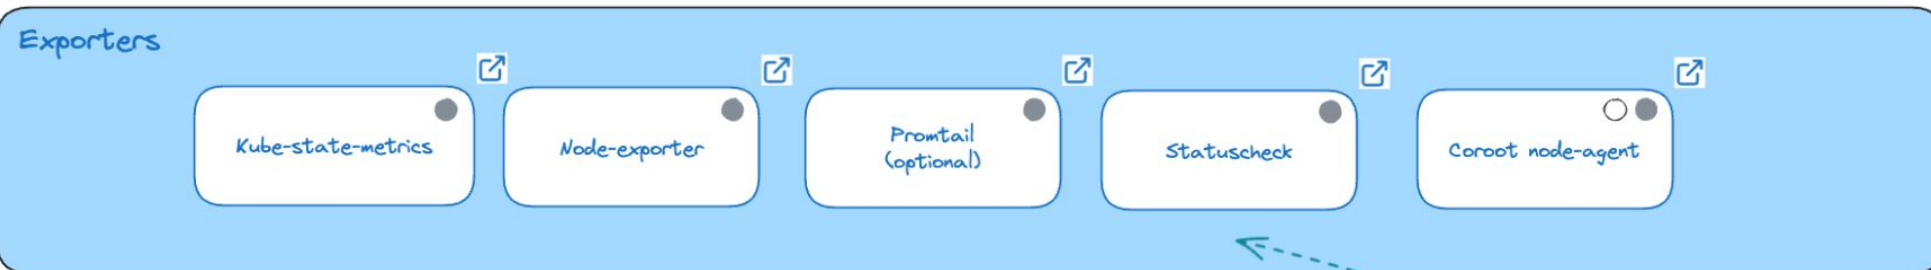
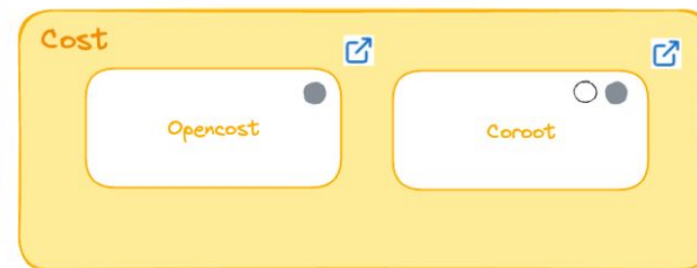
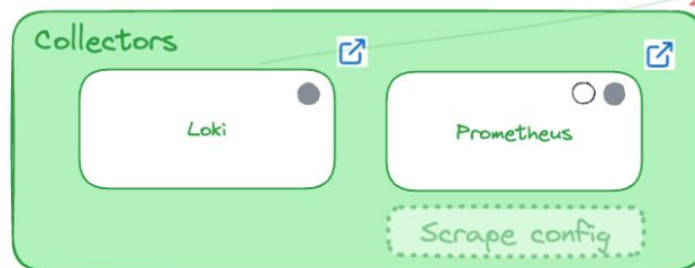
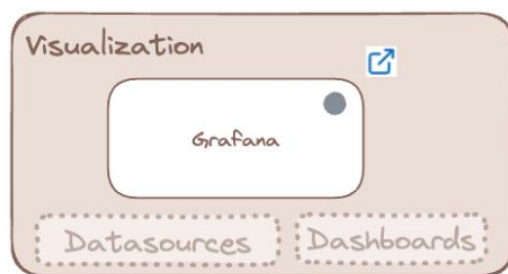
- System tables
- Built-in metrics
- Built-in tracing
- Built-in logs
- Altinity Operator
- Grafana Dashboards

System Tables

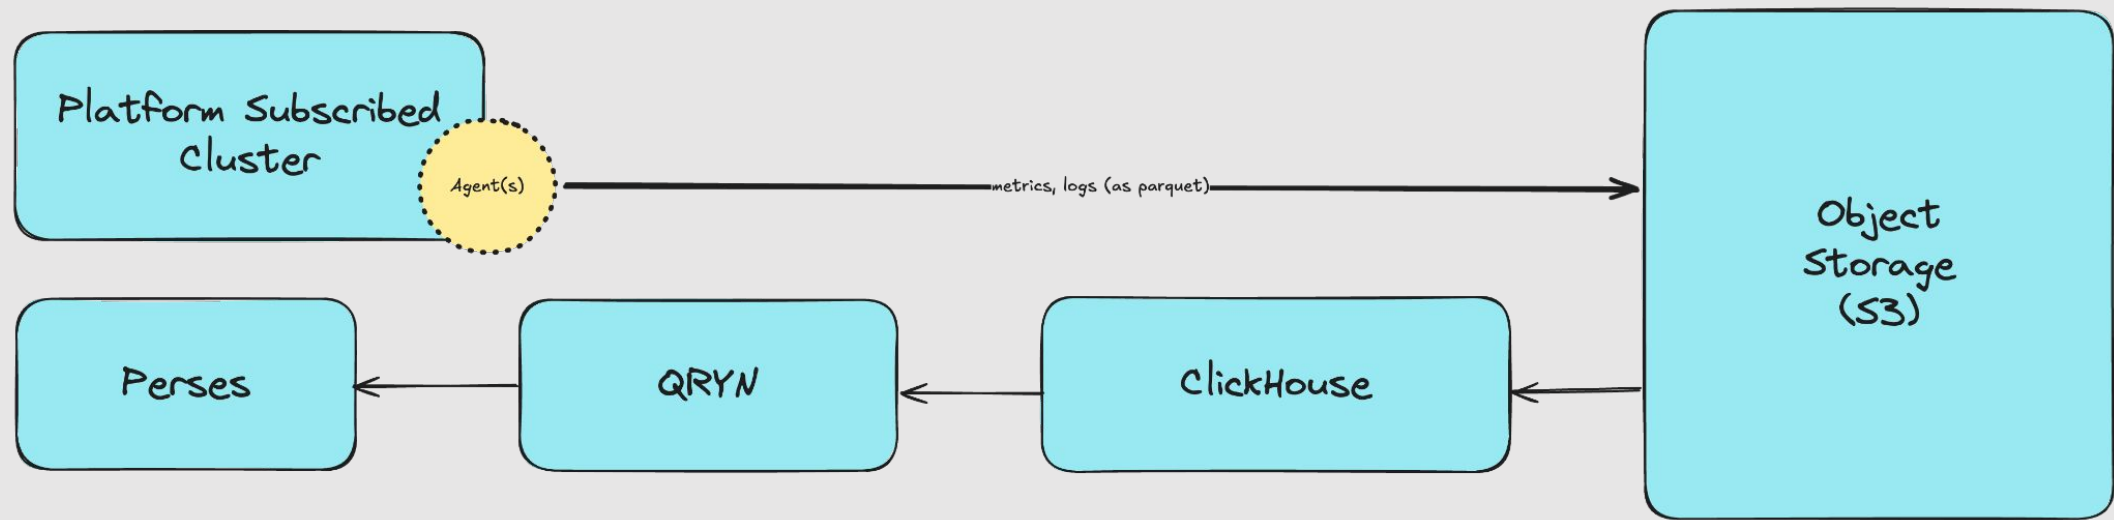
name	database	parts	total_rows
asynchronous_metric_log	system	84	31806721
metric_log	system	78	40947
part_log	system	81	146948
query_log	system	133	74777
query_views_log	system	133	8063
trace_log	system	85	1891241



Observability



namespace=altinity-cloud-system




▼  grafana

▼  anywhere-dashboards

 anywhere-cloud-connect.json


 anywhere-etcd.json


 anywhere.json

▼  dashboards

 acm-metrics-renderer.json

 anywhere-clickhouse-backup.json

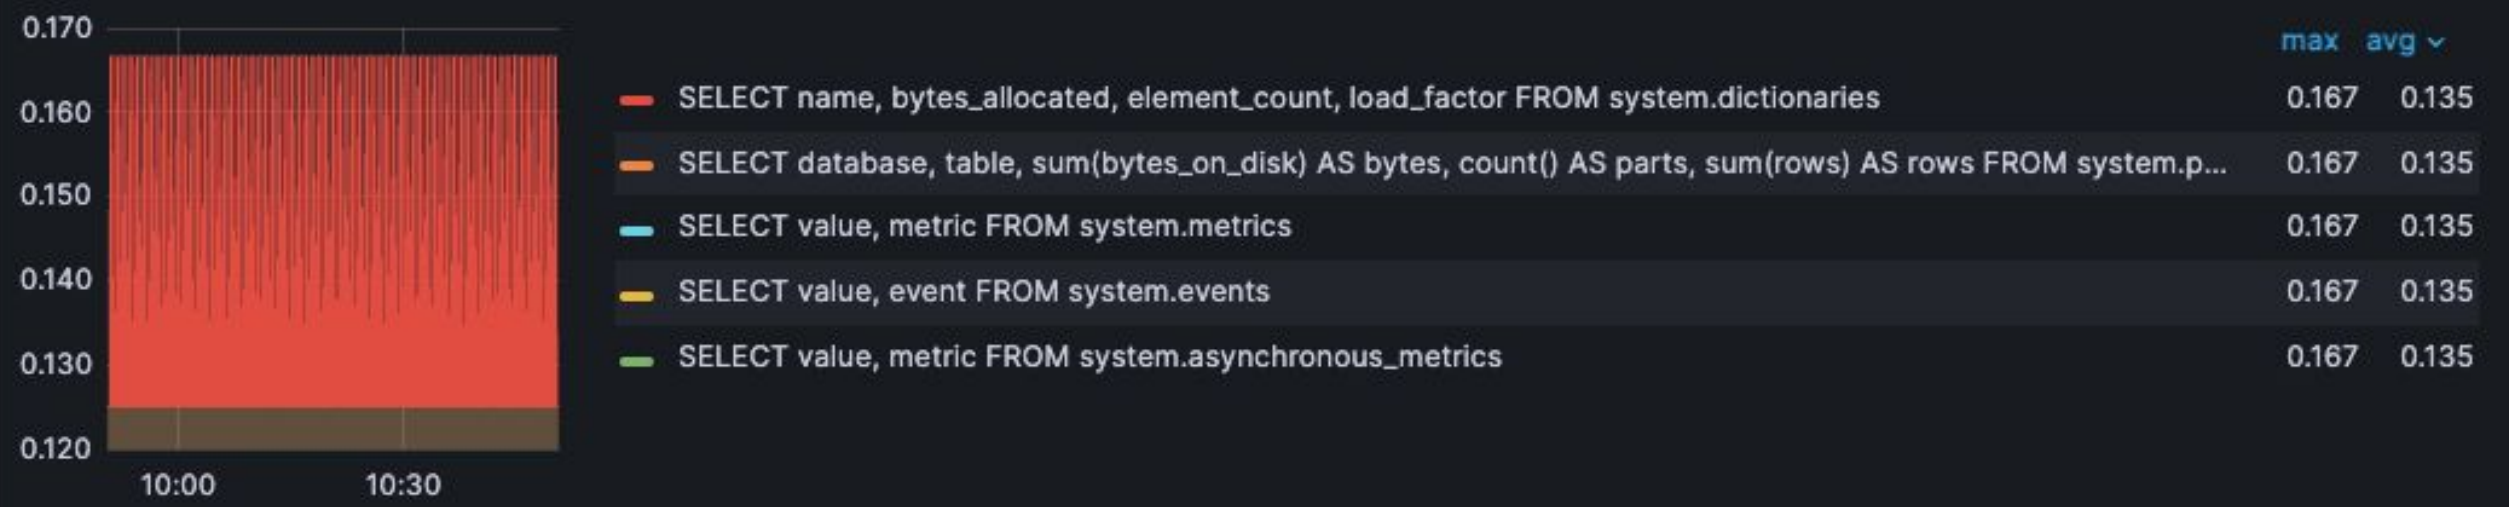
 anywhere-clickhouse-kafka.json

 anywhere-clickhouse-operator.json

Dashboards

Top charts

Top 5 request's rate by type: All; user: All; query kind: All



Top slow queries by type: All; user: All; qu...

query	duration_ms
SELECT name, bytes_allocated, element_count, load_factor FROM system.dictionaries	0.08
SELECT value, metric FROM system.metrics	0.21
SELECT value, metric FROM	

Top memory consumers by type: All; user:...

query	usage	count
system.dictionaries		▾
SELECT version()	8.95 KiB	480
SELECT database, table, sum(bytes_on_disk) AS bytes, count() AS parts,		

Top failed queries by user: All; query kind:...


No data to show ⓘ

Alerting

Alerts for known troublemakers

CrashLoopBackOff
KubeJobFailed
KubeDaemonSetRolloutStuck
KubePersistentVolumeUsage
PrometheusIsDown

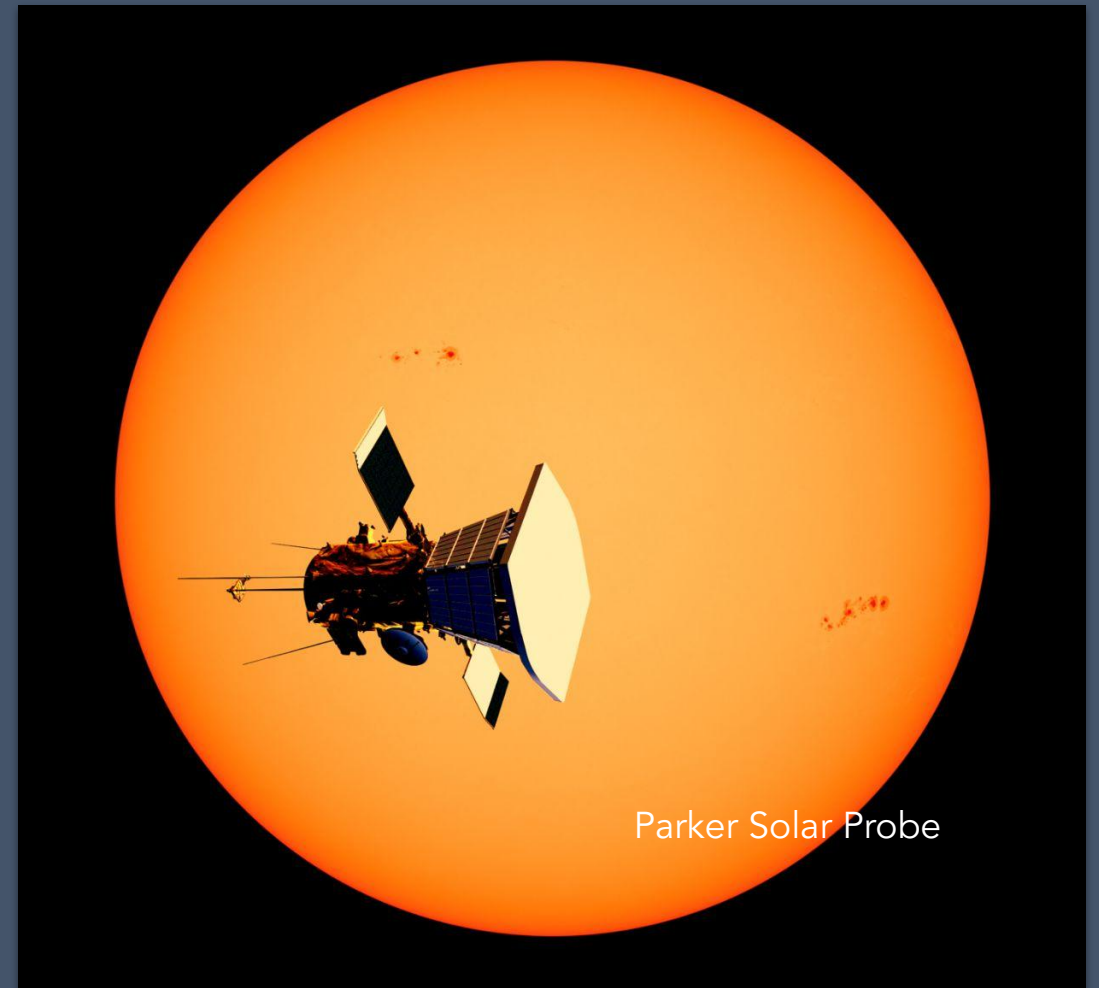
A complete observability solution



eBPF

Provides external observability into any syscalls made by a target process.

Allows network request mapping.



coroot:~#

default ▾

🔍 search for apps and nodes



🕒 last hour ▾



Overview

HEALTH

SERVICE MAP

TRACES

NODES

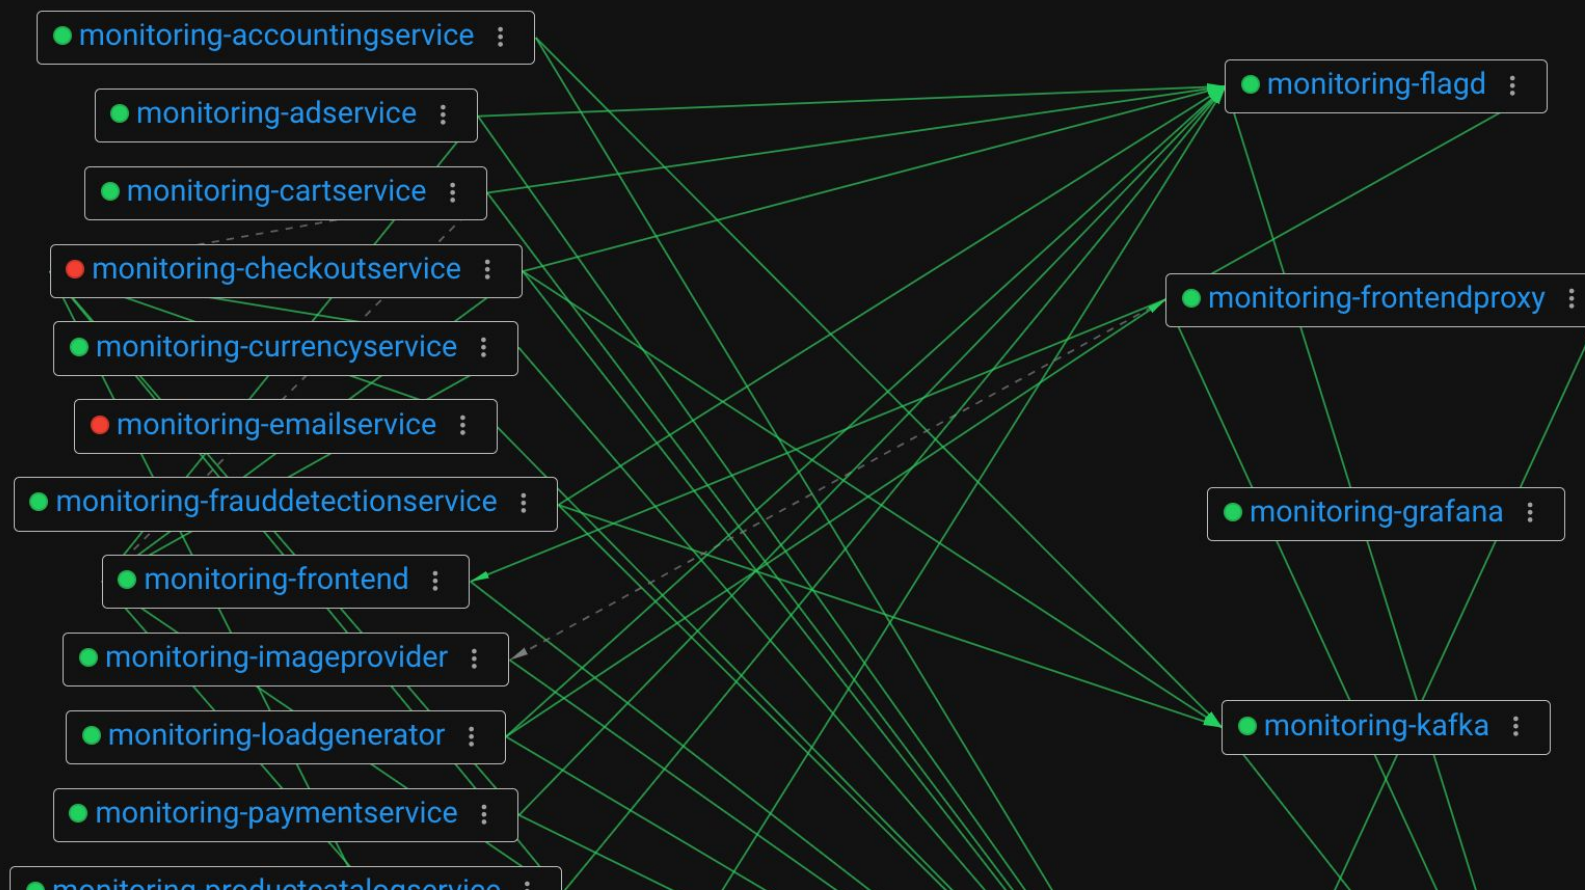
DEPLOYMENTS

COSTS

🔍 search

namespaces

monitoring ✕ ▾

☒ application ☐ control-plane ☒ monitoring +



#1 PRODUCT OF THE MONTH

User Experience



Cloud

Blog

Get Started

Introduction

Supported APIs

✓ Installation

Settings

✓ Data Ingestion

Logs

Metrics

Telemetry

▼ Data Ingestion

Ingesting data with **qryn** is easy and painless thanks to our **polyglot** design.

Use any **Agent** or **Library** compatible with *Opentelemetry*, *Loki*, *Datadog*, *Elastic*, *Prometheus*, *Tempo*, *Graphite*, *Pyroscope* & more

Logs

Metrics

Telemetry

Profiling

Custom

Get started using the [Log Ingestion](#) section



▼ API Support

Food for Thought...

1. Continuity
2. Annotations, Normalization, Filtering (ETL to ELT)
3. Combining App & Infra Views
4. Advanced Analytics
5. Data-mining + Model Training

Wrapping Up

Start with what you have

Focus on known troublemakers

Guard against axes of change

Never stop adapting

Thank You OpenSearch Community!



Connect with me



We're seeking OSS contributors!
Join our Slack
altinity.com/slack