



A Roundup of Observability Datastores

Josh Lee • Altinity • OSOD • Oct. 23, 2025

**"A Developer will never ask you,
'Hey, what filesystem is that?'"**

— Patrick McFadin



Josh Lee

Open Source Advocate
Altinity

*ClickHouse® is a registered trademark of ClickHouse, Inc.
Altinity is not affiliated with or associated with ClickHouse, Inc.
We are but humble open source contributors*

**Observability = Visibility +
Understanding**

50x

Observability data vs system data

What are we storing?

Metrics, Traces, Logs, Profiles, Events

Labels/Tags

Resource Metadata

Graphs & Topologies

Snapshots & Deltas

Configuration (e.g. alerts, users, dashboards)

What do we need for observability?

Fast streaming writes

Efficient compression & storage

Time-oriented management

“Real-time” analytics

"Anything you can do with a group by, that's what analytics is"

—Peter Marshall

More Requirements

Fast multi-row analytics

Full-text search

Tag/label search

Fast, frequent "last point" reads

Updates?

Database Archetypes

OLTP

OLAP

TSDB

Search/Analytics

Introducing the cast of characters

Postgres (OLTP)

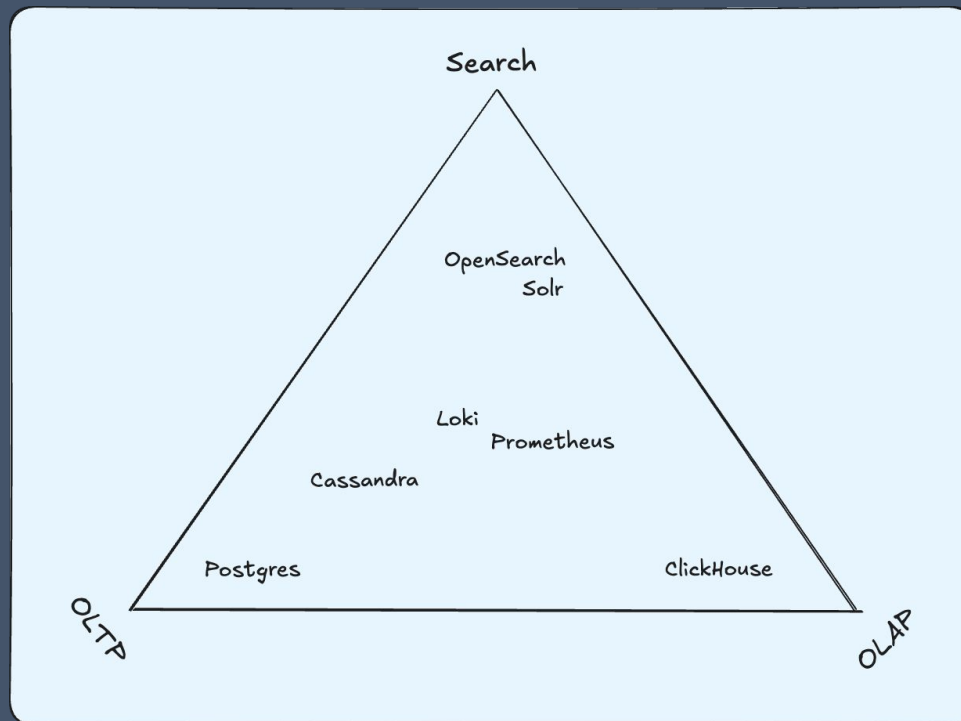
Cassandra (OLTP)

OpenSearch (Search & Analytics)

Prometheus (TSDB)

ClickHouse (OLAP)

Taxonomies are challenging



Storage on disk

Database Storage Styles

Heap Pages + Commit Log

Time-series Blocks






Parts / Segments

Heap Pages

* the JBOD of storage styles





Heap Pages

Page marked for deletion








Heap Pages

Vacuum Process removes Page



Heap Pages


Latest page inserted anywhere it fits



TSDB Blocks

Append-only

TSDB Blocks




Immutable Parts / Segments



w/ Background Compaction

Immutable Parts / Segments

Batch insert creates a new part



Parts sorted and merged in background
creating larger parts



Writing Data


Write Ahead Log (WAL) / Commit Log

Buffered, unordered writes stored on disk

Concurrency Control Strategies

- MVCC + Vacuum
- "Tombstone" deletes
- Last-write wins

Balanced Trees (B-Trees)



Now we can build a Postgres

WAL

Heap Pages + MVCC

B-Tree Indexes

Postgres/MySQL/etc.


Optimized for updates/upserts and row-level reads

Strong ACID guarantees

Scaling horizontally is challenging

Analytics & Search Architecture

Log-Structured Merge Tree



and periodically flushed to disk...forming a set of small, sorted files.

Lucene Family:

**Cassandra, Elastic/OpenSearch,
Apache Solr**

A Lucene Query

```
(title:"database systems" OR content:(postgres OR "clickhouse"))  
AND timestamp:[2025-01-01 TO 2025-12-31]  
AND NOT tags:deprecated
```


Cassandra

Wide-event Scalable OLTP

Vector Engines & Search

Inverted Indexes

Bloom Filters

Approximate Nearest Neighbor (ANN Graph)

Inverted Indexes

"cat" → [doc1, doc3, doc7]

"dog" → [doc2, doc5]

"parrot" → [doc1, doc4]

Bloom Filters

No false negatives


"ae7c" → [doc1, doc3, doc7]

"f9b2" → [doc2, doc5]

"8c93" → [doc1, doc4]

Approximate Nearest Neighbor (ANN)

A way to organize and filter vectors




Prometheus (& friends)

Time-series Database

TSDB: Data is naturally ordered by time


Excellent for frequent reads of last-sample



TSDB

No. of time series = cardinality^{dimensionality}


Cardinality Explosions



Row-oriented vs column-oriented storage

Row-oriented Storage


Read all columns in row



Rows compressed minimally or not at all

Column-oriented

Read only selected columns



Columns highly compressed

**59 GB
(100%)**

Read 109
columns

Read 3 columns
from 109

**1.7 GB
(3%)**

Read 3
compressed
columns

21 MB (.035%)


Read 3
compressed
columns over
8 threads

**2.6 MB
(.0044%)**

ClickHouse

Column-oriented MergeTree

Unmerged,
freshly
inserted
part




Fully
merged
part




Query efficiency

Sparse Indexes

Quickly & cheaply find data
based on an ordered key



Which to choose?



VictoriaMetrics

TSDB meets MergeTree

Loki

Uncomplicated logging
(with label indexes)

Honorable Mentions

Cortex, Thanos, Mimir, TimecaleDB, Solr, Druid...

At (very) small scale

Just use what you have until it breaks (Postgres)

Hooked-on full-text search

OpenSearch has your back

One database for everything

ClickHouse is pretty cool

Wide-event analytics

ClickHouse is awesome

Filtering heavily before analyzing

OpenSearch is also a good choice here


**Lots of
"last-sample" reads + alerts**

Choose a TSDB like Prometheus or VictoriaMetrics

Wide-events analytics with transactional guarantees

Cassandra or Postgres->ClickHouse

Database (Orientation)	Style/QL	Storage	Indexes	Use Case
Postgres (Row)	OLTP/SQL	Heap Pages	B-Tree	Update/Upsert with Guarantees
Cassandra (Document)	OLTP/CQL	Lucene Segments	Inverted	Scalable Upserts
Prometheus (Columnar*)	TSDB/PromQL	TSDB files	By label	Time-series metrics, alerting
OpenSearch (Document)	Search/LuceneQL	Lucene Segments	Inverted, Bloom Filter, ANN	Full-text search, analytics
ClickHouse (Columnar)	OLAP/SQL	MergeTree Parts	Sparse, Inverted, and more...	Wide-event analytics

A full-page background image of a surfer riding a large, curling blue wave. The surfer is positioned in the lower right corner, wearing a dark wetsuit and riding a light-colored surfboard. The wave is a deep blue with white foam at the crest. The text is overlaid on the left side of the image.

Thank you and happy querying!

Josh Lee - Altinity