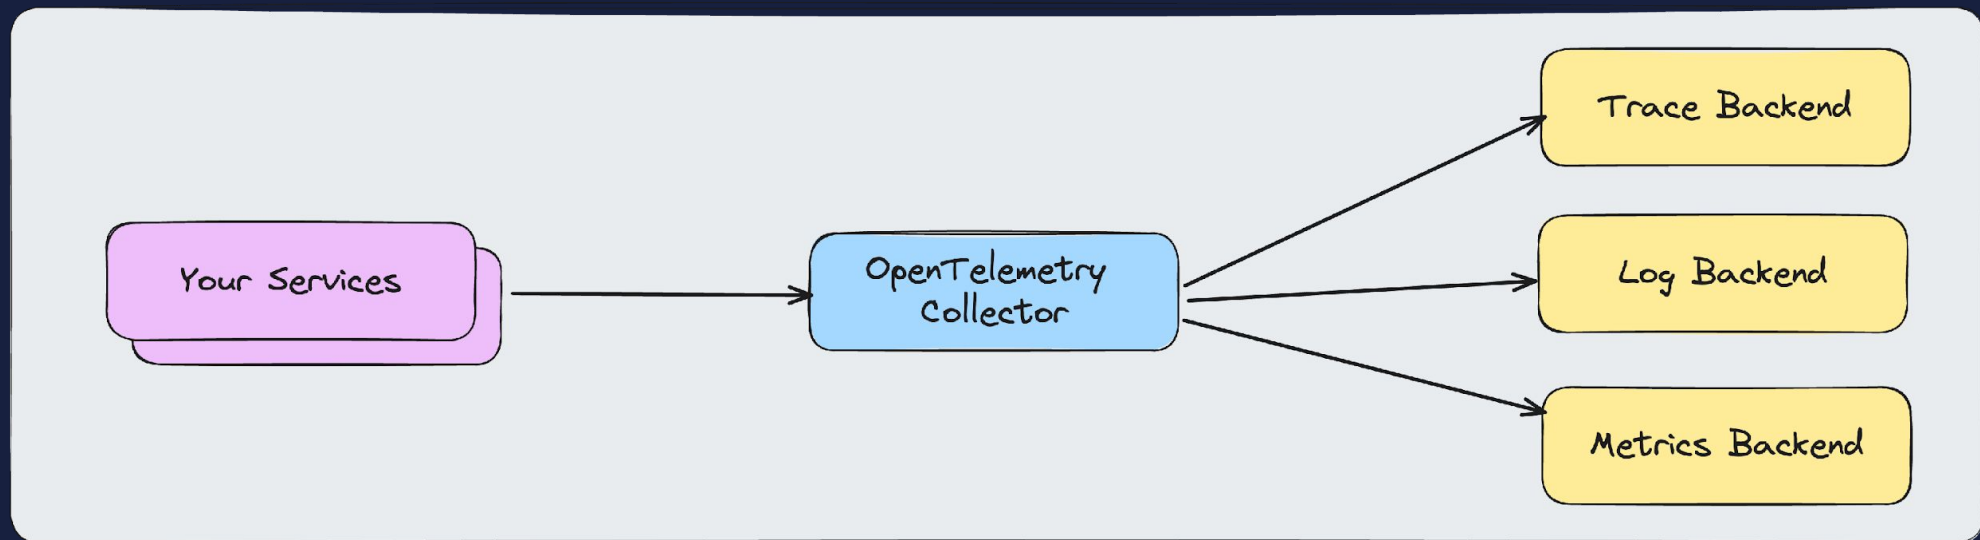


# O11y in One: ClickHouse® as a Unified Telemetry Database

# How I usually start...




"The OpenTelemetry project does not include any kind of database or backend UI."

# 6

Minimum no. of o11y tools deployed by a  
typical organization

— Grafana State of Observability Report, 2023

# What we really need...





Josh Lee  
Open Source Advocate  
*Altinity*

*Altinity® is a Registered Trademark of Altinity, Inc. ClickHouse® is a registered trademark of ClickHouse, Inc.;  
Altinity is not affiliated with or associated with ClickHouse, Inc.  
We are but humble open source contributors*

# Agenda

- 01 — What's the problem with disparate systems?
- 02 — ClickHouse for Observability
- 03 — Full-stack Solutions





# Challenges with Disparate Telemetry Systems





# What are we storing?

Metrics, traces, logs, profiles, events

Resource metadata

Graphs & topologies

Snapshots & deltas

Configuration

# Is There a Silver Bullet?

Real-time analytics

Efficient compression

Full-text search

Relational

Petabyte-scale

# No. Obviously.

... but ClickHouse comes pretty close.

# Introducing ClickHouse

- SQL-compatible
- Massively scaleable
- Really, really fast

# ClickHouse for Observability




# Telemetry is WORM


Write-Once, Read-Many

# B-Trees: Optimized for Reads

# Log-Structured Merge Trees: Optimized for ingestion




# Log-Structured Merge Trees: Background compaction




Compaction continues creating fewer, larger and larger files


Part



Part

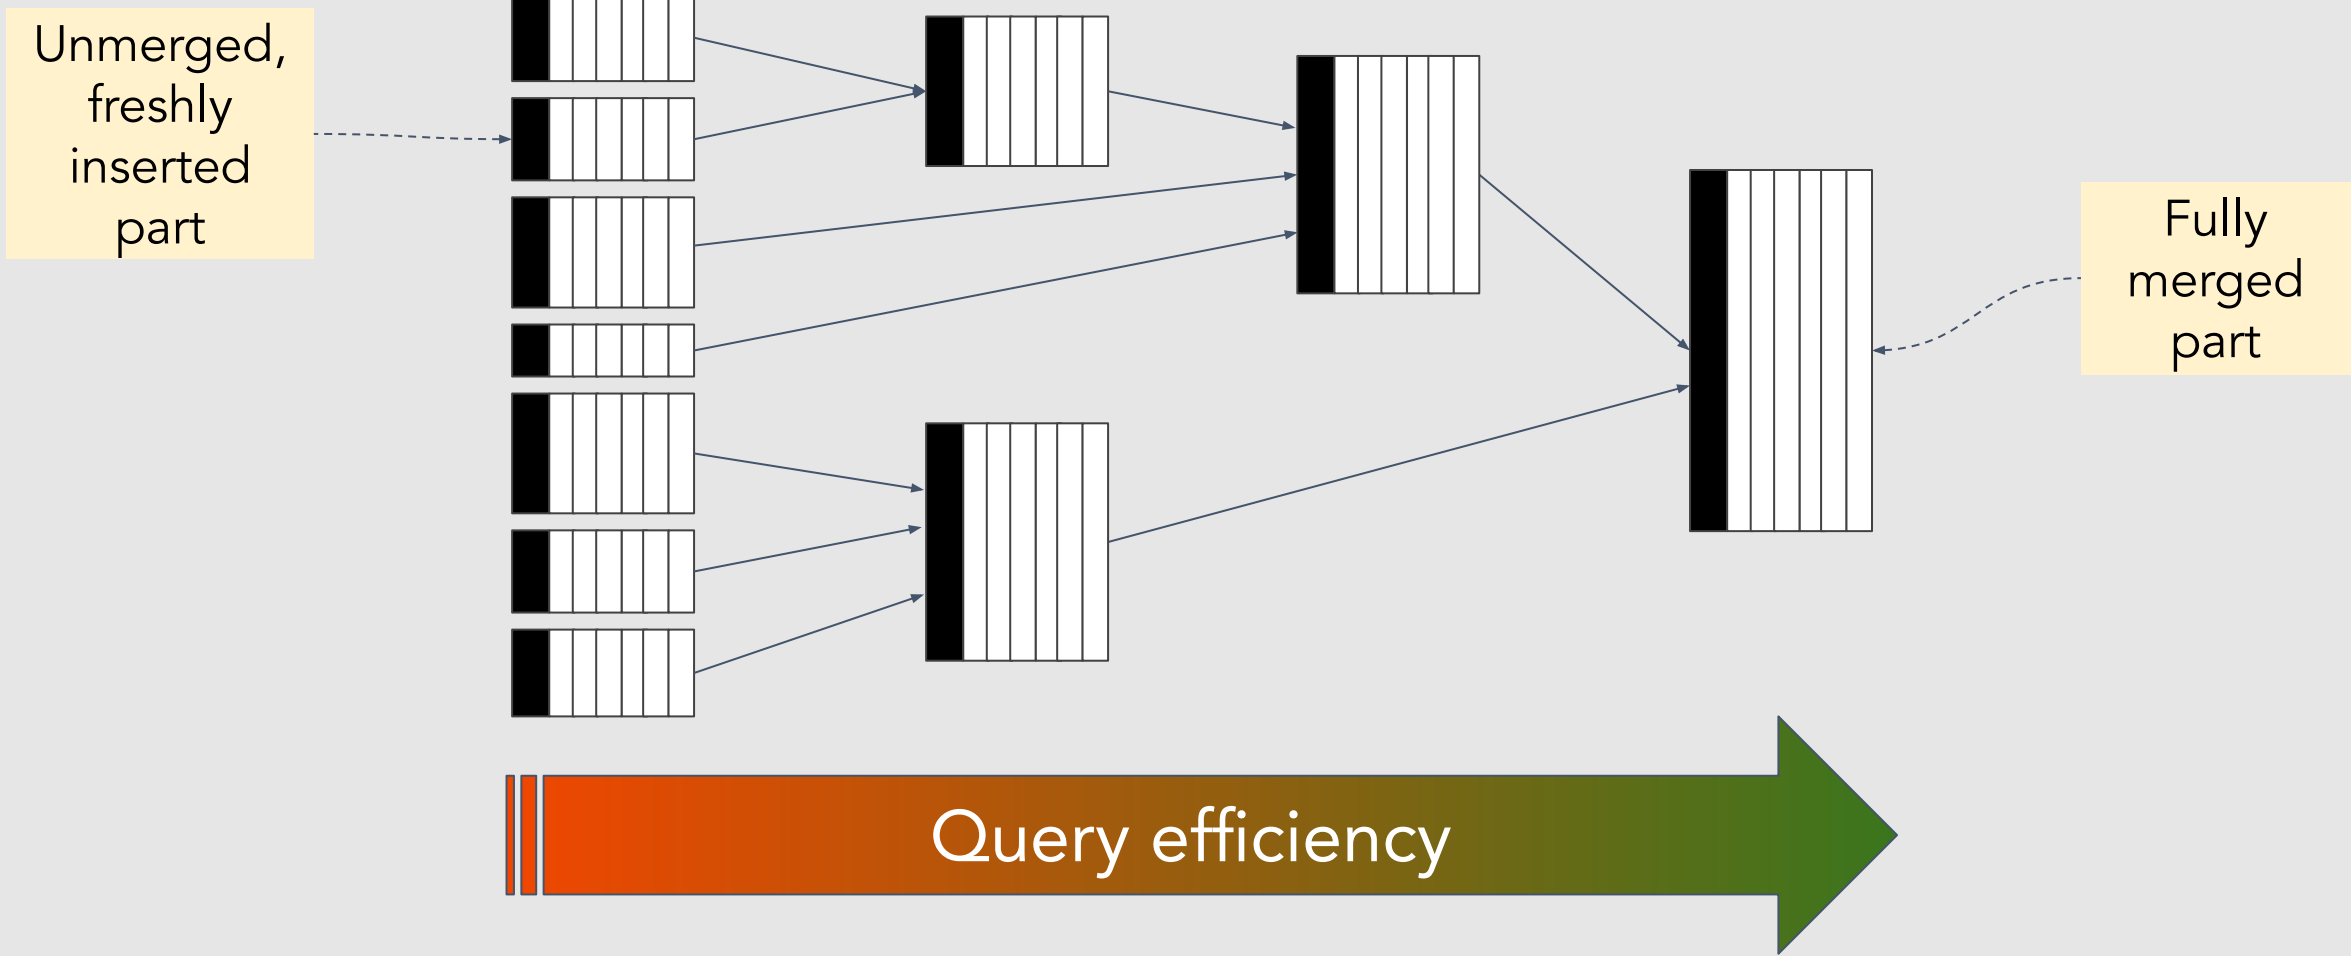


Rewritten, Bigger Part



Update and delete also rewrite parts





ClickHouse for Observability

# How does this help?

- Fast writes
- Time-friendly
- Easy cleanup
- Cost-effective

# Data Transformation & Management

- Materialized Views
- TTL
- Tiered storage

ClickHouse for Observability

# Integrations

- Grafana Datasource Plugin
- Jaeger w/ ClickHouse backend
- cLoki
- Kafka table engine

ClickHouse for Observability


# New & Upcoming

- Time Series Table Engine
- PromQL Support
- Awesome compression for JSON columns / arbitrary schema



# ClickHouse for Observability

## Integrations via OpenTelemetry



ClickHouse for Observability

# More Benefits

- Excellent compression, even with variable schemas
- Practically unlimited cardinality
- Horizontally scalable ingestion & querying

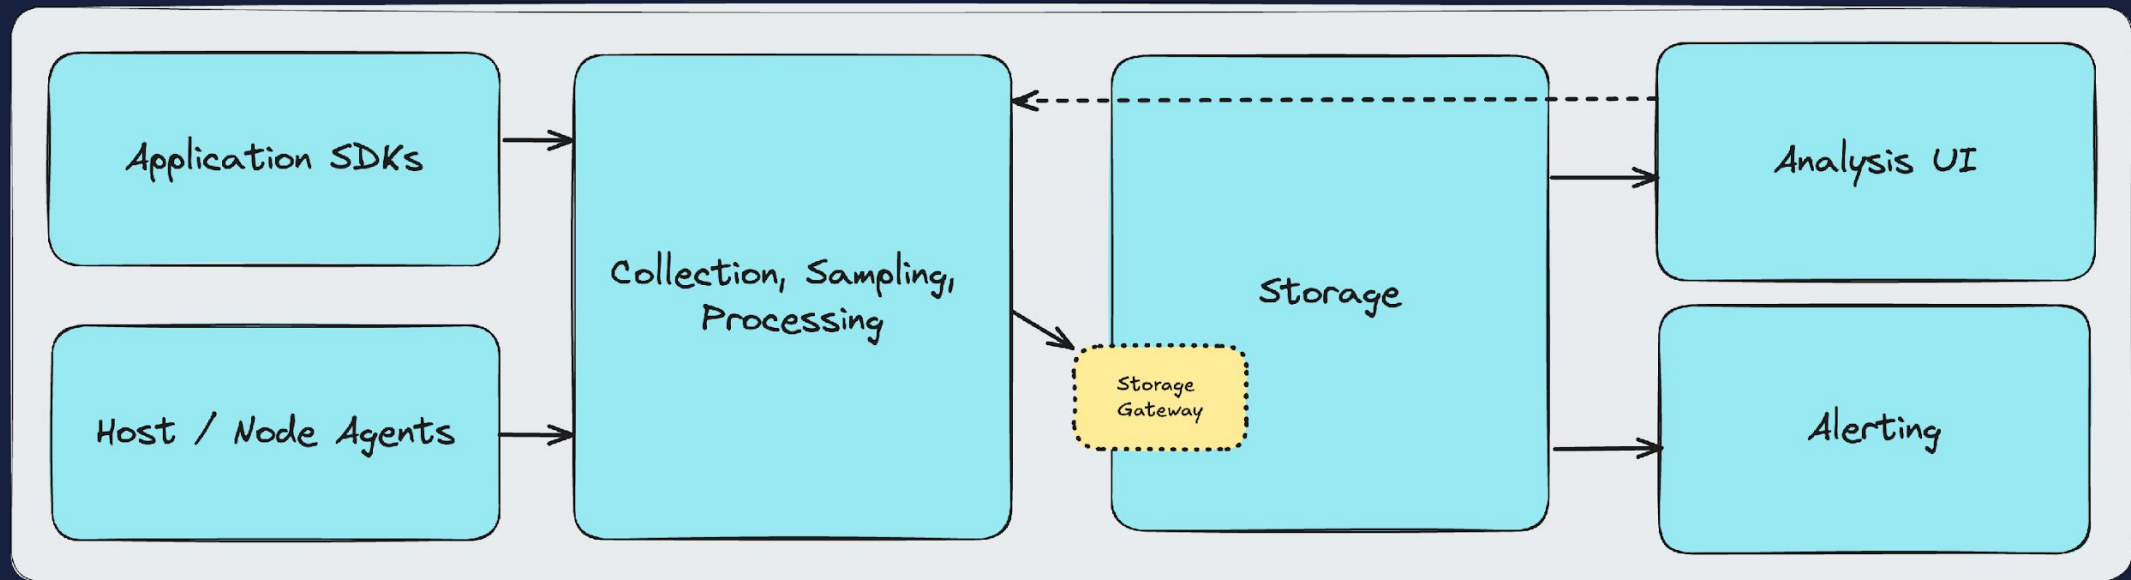
ClickHouse for Observability

# Challenges

- SQL is not PromQL\*
- Overly complex for small data volumes\*
- Not a turn-key solution

"The OpenTelemetry project does not include any kind of database or *backend UI*."

# We need a complete observability solution





SigNoz

Coroot

qryn

HyperDX / ClickStack

DIY

# Coroot

Batteries-included, no-code  
observability

coroot:~#

default ▾

🔍 search for apps and nodes



last hour ▾



## Overview

HEALTH

SERVICE MAP

TRACES

NODES

DEPLOYMENTS

COSTS

🔍 search

namespaces

monitoring X ▾

☒ application ☐ control-plane ☒ monitoring +

2 SLO violation 2 Warning 3 Errors in logs 0 Integration required 19 OK

Application	Type	Errors	Latency	Upstreams	Instances	Restarts	CPU	Mem	I/O load	Disk	Net	DNS	Logs
monitoring-checkoutservice	golang	–	10s	8/9	1/1	–	shortage	–	–	–	0.8ms	–	–
monitoring-emailservice		–	10s	1/1	1/1	–	shortage	–	–	–	<0.1ms	–	–
chi-monitoringdb-monitoring-0-0	clickhouse	–	2s	–	1/1	–	–	–	0.004	100%	–	–	14 unique errors
monitoring-prometheus-server	prometheus	–	633ms	1/1	1/1	–	–	–	–	–	failed conns	–	–
monitoring-frontend		<1%	10s	6/7	1/1	–	–	–	–	–	2ms	–	2 unique errors
monitoring-opentelemetry-collector-...	golang	–	4s	1/1	2/2	–	–	–	–	–	0.8ms	–	4 unique errors
monitoring-otelcol	golang	–	602ms	4/4	1/1	–	–	–	–	–	<0.1ms	–	1 unique error
monitoring-accountingservice		–	–	2/2	1/1	–	–	–	–	–	0.9ms	–	–
monitoring-adservice	java	–	27ms	2/2	1/1	–	–	–	–	–	<0.1ms	–	–
monitoring-cartservice	dotnet	–	249ms	3/3	1/1	–	–	–	–	–	0.8ms	–	–
monitoring-currencyservice		–	3s	1/1	1/1	–	–	–	–	–	0.7ms	–	–
monitoring-flagd	golang	–	10s	1/1	1/1	–	–	–	–	–	<0.1ms	–	–
monitoring-frauddetectionservice	java	–	–	3/3	1/1	–	–	–	–	–	0.1ms	–	–

coroot:~#

default ▾

🔍 search for apps and nodes



🕒 last hour ▾



## Overview

HEALTH

**SERVICE MAP**

TRACES

NODES


DEPLOYMENTS

COSTS

🔍 search

namespaces

monitoring ✕ ▾

☒ application ☐ control-plane ☒ monitoring +

coroot:~#

default ▾

🔍 search for apps and nodes



🕒 last hour ▾



## Applications / monitoring-emailservice

● monitoring-checkoutservice  
ns:monitoring

● monitoring-emailservice :  
ns:monitoring

monitoring-emailservice-75d6f...

● monitoring-otelcol  
ns:monitoring

● SLO

● INSTANCES

● CPU

● MEMORY

● NET

● LOGS

PROFILING

TRACING

● Errors: 1 error occurred

Condition: the number of messages with the ERROR and CRITICAL severity levels > 0

● Using container logs ([configure](#))

Filter:

🔍 Filter messages

Query

View:

☰ Messages

✦ Patterns

↑ Newest first

↓ Oldest first


Limit: 100



# ClickStack

Bundled OTel + ClickHouse +  
HyperDX (visualization and query  
UI)

Select Inserts Infrastructure





# coroot

- eBPF-based Node-Agent
- OTLP ingestion via Collector Gateway
- Uses (mostly) standard OpenTelemetry Exporter schema + new schema for profiles

# clickstack

- OTel-collector based agent
- Uses standard OTel-exporter schema
- Schema-Agnostic Visualization UI
- More info @ 3PM

# qryn

- Uses on its own collector exporter / collector distribution
- Exposes Tempo, Loki, OTLP, and Prometheus APIs
- Projects into compatible formats using Materialized Views



# Schema Considerations

# Schema Considerations

- ZSTD Compression
- Delta encoding
- Bloom filter indexes for maps (resources) and logs
- MergeTree, partitioned on time
- 7-day TTL

# OpenTelemetry Collector Exporter for ClickHouse


- Maps for metadata
- Efficient full-body text-search
- Materialized View for span durations



- Fingerprints for unique time series
- Indexed labels (via Materialized Views)
- Allows for efficient updates (ReplacingMergeTree)
- Null Engine for raw ingest

# Scaling for Production

# Managing Multiple Collectors



# The Altinity Operator

- PVC management
- Rolling upgrades
- Built-in monitoring

# Alerting & Other Considerations



## Conclusion


# Why Unified Observability Storage?

Easier onboarding  
Simplified management

Simplified scaling  
Cost management

Cross-signal correlation with shared  
metadata

# Thank You



London Meetup Oct  
23rd - Open Lakehouses



My LinkedIn  
(will post slides)